

## Deep Learning Libraries

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
pip install numpy opencv-python tensorflow matplotlib seaborn scikit-learn
```

Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (1.26.4)

Requirement already satisfied: opencv-python in /usr/local/lib/python3.10/dist-packages (4.10.0.84)

Requirement already satisfied: tensorflow in /usr/local/lib/python3.10/dist-packages (2.17.1)

Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (3.8.0)

Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-packages (0.13.2)

Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (1.5.2)

Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.4.0)

Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.6.3)

Requirement already satisfied: flatbuffers>=24.3.25 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (24.3.25)

Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.6.0)

Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.2.0)

Requirement already satisfied: h5py>=3.10.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.12.1)

Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (18.1.1)

Requirement already satisfied: ml-dtypes<0.5.0,>=0.3.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.4.1)

Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.4.0)

Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from tensorflow) (24.2)

Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<5.0.0dev,>=3.20.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (4.25.5)

Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.32.3)

Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from tensorflow) (75.1.0)

Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.16.0)

Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.5.0)

Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (4.12.2)

Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.16.0)

Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.68.0)

Requirement already satisfied: tensorboard<2.18,>=2.17 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.17.1)

Requirement already satisfied: keras>=3.2.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.5.0)

Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.37.1)

Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.3.1)

Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (4.55.0)

Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.4.7)

Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (11.0.0)

Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (3.2.0)

Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (2.8.2)

Requirement already satisfied: pandas>=1.2 in /usr/local/lib/python3.10/dist-packages (from seaborn) (2.2.2)

Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.13.1)

Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.4.2)

Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (3.5.0)

Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.10/dist-packages (from astunparse>=1.6.0->tensorflow) (0.45.0)

Requirement already satisfied: rich in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->tensorflow) (13.9.4)

Requirement already satisfied: namex in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->tensorflow) (0.0.8)

Requirement already satisfied: optree in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->tensorflow) (0.13.1)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.2->seaborn) (2024.2)

Requirement already satisfied: tzdata>=2022.7 in

/usr/local/lib/python3.10/dist-packages (from pandas>=1.2->seaborn) (2024.2)  
Requirement already satisfied: charset-normalizer<4,>=2 in  
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0-  
>tensorflow) (3.4.0)  
Requirement already satisfied: idna<4,>=2.5 in  
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0-  
>tensorflow) (3.10)  
Requirement already satisfied: urllib3<3,>=1.21.1 in  
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0-  
>tensorflow) (2.2.3)  
Requirement already satisfied: certifi>=2017.4.17 in  
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0-  
>tensorflow) (2024.8.30)  
Requirement already satisfied: markdown>=2.6.8 in  
/usr/local/lib/python3.10/dist-packages (from tensorboard<2.18,>=2.17-  
>tensorflow) (3.7)  
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in  
/usr/local/lib/python3.10/dist-packages (from tensorboard<2.18,>=2.17-  
>tensorflow) (0.7.2)  
Requirement already satisfied: werkzeug>=1.0.1 in  
/usr/local/lib/python3.10/dist-packages (from tensorboard<2.18,>=2.17-  
>tensorflow) (3.1.3)  
Requirement already satisfied: MarkupSafe>=2.1.1 in  
/usr/local/lib/python3.10/dist-packages (from werkzeug>=1.0.1-  
>tensorboard<2.18,>=2.17->tensorflow) (3.0.2)  
Requirement already satisfied: markdown-it-py>=2.2.0 in  
/usr/local/lib/python3.10/dist-packages (from rich->keras>=3.2.0->tensorflow)  
(3.0.0)  
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in  
/usr/local/lib/python3.10/dist-packages (from rich->keras>=3.2.0->tensorflow)  
(2.18.0)  
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.10/dist-  
packages (from markdown-it-py>=2.2.0->rich->keras>=3.2.0->tensorflow) (0.1.2)

```
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import zipfile
import cv2
```

*# Deep Learning Libraries*

```
import tensorflow as tf
from tensorflow.keras import layers, models, optimizers, regularizers
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.applications import VGG16, ResNet50
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
```

## 1. Load the images.

## 2. Discard images that correspond to ages outside the range [10, 90].

```
def load_and_preprocess_data(zip_file_path):
    with zipfile.ZipFile(zip_file_path, 'r') as zip_ref:
        zip_ref.extractall('dataset')

    images = []
    ages = []

    dataset_folder = "/content/drive/MyDrive/FCAI /utkcropped"

    # Walk through extracted directory
    for filename in os.listdir(dataset_folder):
        if filename.endswith('.jpg'):
            parts = filename.split('_')
            try:
                age = int(parts[0])

                if 10 <= age <= 90:
                    img_path = os.path.join(dataset_folder, filename)
                    img = cv2.imread(img_path)
                    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
                    img = cv2.resize(img, (64, 64))

                    images.append(img)
                    ages.append(age)
            except ValueError:
                continue

    X = np.array(images) / 255.0
    y = np.array(ages)

    return X, y
```

## Pre Model

```
# Split the data
def split_data(X, y, test_size=0.2, val_size=0.2):
    from sklearn.model_selection import train_test_split

    X_train_val, X_test, y_train_val, y_test = train_test_split(
        X, y, test_size=test_size, random_state=42
    )

    X_train, X_val, y_train, y_val = train_test_split(
        X_train_val, y_train_val, test_size=val_size, random_state=42
    )
```

```

    return X_train, X_val, X_test, y_train, y_val, y_test

def train_model(model, X_train, y_train, X_val, y_val,
data_augmentation=None):
    # Callbacks
    early_stopping = EarlyStopping(
        monitor='val_loss',
        patience=5,
        restore_best_weights=True
    )

    reduce_lr = ReduceLROnPlateau(
        monitor='val_loss',
        factor=0.2,
        patience=3,
        min_lr=0.00001
    )

    # Training
    if data_augmentation:
        history = model.fit(
            data_augmentation.flow(X_train, y_train, batch_size=32),
            validation_data=(X_val, y_val),
            epochs=10,
            callbacks=[early_stopping, reduce_lr]
        )
    else:
        history = model.fit(
            X_train, y_train,
            validation_data=(X_val, y_val),
            epochs=20,
            batch_size=32,
            callbacks=[early_stopping, reduce_lr]
        )

    return history

def plot_training_history(histories, model_names):
    plt.figure(figsize=(15, 5))

    # Training Loss
    plt.subplot(1, 2, 1)
    for history, name in zip(histories, model_names):
        plt.plot(history.history['loss'], label=f'{name} Training Loss')
        plt.plot(history.history['val_loss'], label=f'{name} Validation
Loss')

    plt.title('Model Training and Validation Loss')
    plt.xlabel('Epoch')

```

```
plt.ylabel('Loss')
plt.legend()

plt.tight_layout()
plt.show()
```

## Two large convolutional neural network (CNN) models with custom architectures.

### Modle 1

*# Custom CNN Model 1: Deep CNN*

```
def create_deep_cnn_model1(input_shape=(64, 64, 3)):
    model = models.Sequential([
        layers.Conv2D(32, (3, 3), activation='relu', input_shape=input_shape,
                      kernel_regularizer=regularizers.l2(0.001)),
        layers.BatchNormalization(),
        layers.MaxPooling2D((2, 2)),

        layers.Conv2D(64, (3, 3), activation='relu',
                      kernel_regularizer=regularizers.l2(0.001)),
        layers.BatchNormalization(),
        layers.MaxPooling2D((2, 2)),

        layers.Conv2D(128, (3, 3), activation='relu',
                      kernel_regularizer=regularizers.l2(0.001)),
        layers.BatchNormalization(),
        layers.MaxPooling2D((2, 2)),

        layers.Flatten(),
        layers.Dense(256, activation='relu',
                    kernel_regularizer=regularizers.l2(0.001)),
        layers.Dropout(0.5),
        layers.Dense(1, activation='linear')
    ])

    model.compile(optimizer='adam', loss='mean_absolute_error',
                  metrics=['mae'])
    return model

def main():
    zip_file_path = "/content/drive/MyDrive/FCAI /DL_Assignment 1.zip"
    X, y = load_and_preprocess_data(zip_file_path)

    X_train, X_val, X_test, y_train, y_val, y_test = split_data(X, y)

    model = create_deep_cnn_model1()
    history = train_model(model, X_train, y_train, X_val, y_val)
```

```

plt.figure(figsize=(10, 5))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Training and Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()

```

```

# Evaluate Model

```

```

test_loss, test_mae = model.evaluate(X_test, y_test)
print(f"Test Mean Absolute Error: {test_mae}")

```

```

main()

```

```

Epoch 1/20

```

```

412/412 ————— 15s 19ms/step - loss: 13.5787 - mae: 12.9609 -
val_loss: 17.2450 - val_mae: 16.6422 - learning_rate: 0.0010

```

```

Epoch 2/20

```

```

412/412 ————— 3s 6ms/step - loss: 10.5029 - mae: 9.8816 -
val_loss: 12.0934 - val_mae: 11.4556 - learning_rate: 0.0010

```

```

Epoch 3/20

```

```

412/412 ————— 3s 6ms/step - loss: 9.6632 - mae: 9.0219 -
val_loss: 11.4303 - val_mae: 10.7726 - learning_rate: 0.0010

```

```

Epoch 4/20

```

```

412/412 ————— 3s 6ms/step - loss: 9.1362 - mae: 8.4724 -
val_loss: 16.3225 - val_mae: 15.6367 - learning_rate: 0.0010

```

```

Epoch 5/20

```

```

412/412 ————— 3s 6ms/step - loss: 8.9382 - mae: 8.2471 -
val_loss: 11.7494 - val_mae: 11.0350 - learning_rate: 0.0010

```

```

Epoch 6/20

```

```

412/412 ————— 3s 6ms/step - loss: 8.5419 - mae: 7.8207 -
val_loss: 12.8146 - val_mae: 12.0667 - learning_rate: 0.0010

```

```

Epoch 7/20

```

```

412/412 ————— 3s 6ms/step - loss: 7.6754 - mae: 6.9324 -
val_loss: 7.6619 - val_mae: 6.9365 - learning_rate: 2.0000e-04

```

```

Epoch 8/20

```

```

412/412 ————— 3s 6ms/step - loss: 7.3332 - mae: 6.6127 -
val_loss: 7.5944 - val_mae: 6.8894 - learning_rate: 2.0000e-04

```

```

Epoch 9/20

```

```

412/412 ————— 3s 6ms/step - loss: 7.2841 - mae: 6.5835 -
val_loss: 8.1233 - val_mae: 7.4364 - learning_rate: 2.0000e-04

```

```

Epoch 10/20

```

```

412/412 ————— 3s 6ms/step - loss: 7.0181 - mae: 6.3352 -
val_loss: 12.5380 - val_mae: 11.8660 - learning_rate: 2.0000e-04

```

```

Epoch 11/20

```

```

412/412 ————— 3s 6ms/step - loss: 6.8614 - mae: 6.1915 -
val_loss: 10.2125 - val_mae: 9.5521 - learning_rate: 2.0000e-04


```

```

Epoch 12/20

```

412/412 ————— 3s 6ms/step - loss: 6.5656 - mae: 5.9057 -  
val\_loss: 7.2734 - val\_mae: 6.6163 - learning\_rate: 4.0000e-05  
Epoch 13/20  
412/412 ————— 3s 6ms/step - loss: 6.4378 - mae: 5.7816 -  
val\_loss: 7.5986 - val\_mae: 6.9452 - learning\_rate: 4.0000e-05  
Epoch 14/20  
412/412 ————— 3s 6ms/step - loss: 6.4425 - mae: 5.7899 -  
val\_loss: 7.2395 - val\_mae: 6.5902 - learning\_rate: 4.0000e-05  
Epoch 15/20  
412/412 ————— 3s 6ms/step - loss: 6.2136 - mae: 5.5649 -  
val\_loss: 7.3120 - val\_mae: 6.6665 - learning\_rate: 4.0000e-05  
Epoch 16/20  
412/412 ————— 3s 6ms/step - loss: 6.2416 - mae: 5.5967 -  
val\_loss: 7.5407 - val\_mae: 6.8985 - learning\_rate: 4.0000e-05  
Epoch 17/20  
412/412 ————— 3s 6ms/step - loss: 6.1372 - mae: 5.4957 -  
val\_loss: 7.6429 - val\_mae: 7.0043 - learning\_rate: 4.0000e-05  
Epoch 18/20  
412/412 ————— 3s 6ms/step - loss: 6.0581 - mae: 5.4195 -  
val\_loss: 7.2341 - val\_mae: 6.5961 - learning\_rate: 1.0000e-05  
Epoch 19/20  
412/412 ————— 3s 6ms/step - loss: 6.0364 - mae: 5.3987 -  
val\_loss: 7.2224 - val\_mae: 6.5854 - learning\_rate: 1.0000e-05  
Epoch 20/20  
412/412 ————— 3s 6ms/step - loss: 6.0754 - mae: 5.4386 -  
val\_loss: 7.2193 - val\_mae: 6.5832 - learning\_rate: 1.0000e-05



129/129 ————— 1s 7ms/step - loss: 7.0537 - mae: 6.4176  
Test Mean Absolute Error: 6.466558933258057



## Modle 2

```
def create_deep_cnn_model2(input_shape=(64, 64, 3)):
    model = models.Sequential([
        layers.Conv2D(64, (3, 3), activation='relu', input_shape=input_shape,
                      kernel_regularizer=regularizers.l2(0.001)),
        layers.BatchNormalization(),
        layers.MaxPooling2D((2, 2)),

        layers.Conv2D(128, (3, 3), activation='relu',
                      kernel_regularizer=regularizers.l2(0.001)),
        layers.BatchNormalization(),
        layers.MaxPooling2D((2, 2)),

        layers.Conv2D(256, (3, 3), activation='relu',
                      kernel_regularizer=regularizers.l2(0.001)),
        layers.BatchNormalization(),
        layers.MaxPooling2D((2, 2)),

        layers.Conv2D(512, (3, 3), activation='relu',
                      kernel_regularizer=regularizers.l2(0.001)),
        layers.BatchNormalization(),
        layers.MaxPooling2D((2, 2)),

        layers.Flatten(),
        layers.Dense(1024, activation='relu',
                    kernel_regularizer=regularizers.l2(0.001)),
        layers.Dropout(0.5),
        layers.Dense(1, activation='linear')
    ])

    model.compile(optimizer='adam', loss='mean_absolute_error',
                  metrics=['mae'])
    return model

def main1():
    zip_file_path = "/content/drive/MyDrive/FCAI /DL_Assignment 1.zip"
    X, y = load_and_preprocess_data(zip_file_path)

    X_train, X_val, X_test, y_train, y_val, y_test = split_data(X, y)

    model = create_deep_cnn_model2()
    history = train_model(model, X_train, y_train, X_val, y_val)

    plt.figure(figsize=(10, 5))
    plt.plot(history.history['loss'], label='Training Loss')
    plt.plot(history.history['val_loss'], label='Validation Loss')
    plt.title('Model Training and Validation Loss')
    plt.xlabel('Epoch')
```

```

plt.ylabel('Loss')
plt.legend()
plt.show()


test_loss, test_mae = model.evaluate(X_test, y_test)
print(f"Test Mean Absolute Error: {test_mae}")

main1()

Epoch 1/20
412/412 ————— 20s 31ms/step - loss: 14.8094 - mae: 12.9078 -
val_loss: 17.4319 - val_mae: 15.8284 - learning_rate: 0.0010
Epoch 2/20
412/412 ————— 6s 14ms/step - loss: 11.5735 - mae: 10.0127 -
val_loss: 12.8345 - val_mae: 11.3713 - learning_rate: 0.0010
Epoch 3/20
412/412 ————— 6s 14ms/step - loss: 10.4723 - mae: 9.0078 -
val_loss: 23.0845 - val_mae: 21.6226 - learning_rate: 0.0010
Epoch 4/20
412/412 ————— 6s 14ms/step - loss: 9.8285 - mae: 8.3586 -
val_loss: 12.6209 - val_mae: 11.1703 - learning_rate: 0.0010
Epoch 5/20
412/412 ————— 6s 14ms/step - loss: 9.6331 - mae: 8.1867 -
val_loss: 15.2241 - val_mae: 13.8023 - learning_rate: 0.0010
Epoch 6/20
412/412 ————— 6s 14ms/step - loss: 9.1898 - mae: 7.7628 -
val_loss: 12.5908 - val_mae: 11.1641 - learning_rate: 0.0010
Epoch 7/20
412/412 ————— 6s 14ms/step - loss: 8.9211 - mae: 7.4893 -
val_loss: 9.6938 - val_mae: 8.2452 - learning_rate: 0.0010
Epoch 8/20
412/412 ————— 6s 14ms/step - loss: 8.7241 - mae: 7.2663 -
val_loss: 13.5302 - val_mae: 12.0495 - learning_rate: 0.0010
Epoch 9/20
412/412 ————— 6s 14ms/step - loss: 8.3225 - mae: 6.8301 -
val_loss: 10.8508 - val_mae: 9.3377 - learning_rate: 0.0010
Epoch 10/20
412/412 ————— 6s 14ms/step - loss: 8.3284 - mae: 6.8008 -
val_loss: 9.5319 - val_mae: 7.9922 - learning_rate: 0.0010
Epoch 11/20
412/412 ————— 6s 14ms/step - loss: 8.2354 - mae: 6.6958 -
val_loss: 8.8288 - val_mae: 7.2620 - learning_rate: 0.0010
Epoch 12/20
412/412 ————— 6s 14ms/step - loss: 8.0376 - mae: 6.4525 -
val_loss: 8.8364 - val_mae: 7.2250 - learning_rate: 0.0010
Epoch 13/20
412/412 ————— 6s 14ms/step - loss: 7.7545 - mae: 6.1349 -
val_loss: 9.0223 - val_mae: 7.3993 - learning_rate: 0.0010
Epoch 14/20
412/412 ————— 6s 14ms/step - loss: 7.6382 - mae: 6.0053 -

```

val\_loss: 10.1613 - val\_mae: 8.5197 - learning\_rate: 0.0010  
 Epoch 15/20  
 412/412 ————— 6s 14ms/step - loss: 6.8295 - mae: 5.2073 -  
 val\_loss: 8.1589 - val\_mae: 6.5971 - learning\_rate: 2.0000e-04  
 Epoch 16/20  
 412/412 ————— 6s 14ms/step - loss: 5.9676 - mae: 4.4239 -  
 val\_loss: 8.1650 - val\_mae: 6.6723 - learning\_rate: 2.0000e-04  
 Epoch 17/20  
 412/412 ————— 6s 14ms/step - loss: 5.6444 - mae: 4.1656 -  
 val\_loss: 7.9789 - val\_mae: 6.5380 - learning\_rate: 2.0000e-04  
 Epoch 18/20  
 412/412 ————— 6s 14ms/step - loss: 5.3248 - mae: 3.8948 -  
 val\_loss: 7.8710 - val\_mae: 6.4712 - learning\_rate: 2.0000e-04  
 Epoch 19/20  
 412/412 ————— 6s 14ms/step - loss: 5.1489 - mae: 3.7571 -  
 val\_loss: 7.7924 - val\_mae: 6.4248 - learning\_rate: 2.0000e-04  
 Epoch 20/20  
 412/412 ————— 6s 14ms/step - loss: 4.9074 - mae: 3.5470 -  
 val\_loss: 7.8498 - val\_mae: 6.5105 - learning\_rate: 2.0000e-04



129/129 ————— 1s 11ms/step - loss: 7.6636 - mae: 6.2959  
 Test Mean Absolute Error: 6.37539005279541

#Two smaller models that apply transfer learning by taking the features extracted by a base CNN model of a well-known architecture as their input, and learning to leverage these features for the age estimation task.

## ResNet50

```
import tensorflow as tf
from tensorflow.keras import layers, models, optimizers, regularizers
```

```

from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.applications import VGG16, ResNet50
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau

```

```

def load_and_preprocess_data(zip_file_path):

```

```

    with zipfile.ZipFile(zip_file_path, 'r') as zip_ref:
        zip_ref.extractall('dataset')

```

```

    images = []
    ages = []

```

```

    dataset_folder = r"/content/drive/MyDrive/FCAI /utkcropped"

```

```

    for filename in os.listdir(dataset_folder):
        if filename.endswith('.jpg'):
            parts = filename.split('_')
            try:
                age = int(parts[0])
                if 10 <= age <= 90:
                    img_path = os.path.join(dataset_folder, filename)
                    img = cv2.imread(img_path)
                    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
                    img = cv2.resize(img, (64, 64))

                    images.append(img)
                    ages.append(age)
            except ValueError:
                continue

```

```

    X = np.array(images) / 255.0
    y = np.array(ages)

```

```

    return X, y

```

```

def create_transfer_learning_model2(input_shape=(64, 64, 3)):
    base_model = ResNet50(weights='imagenet', include_top=False,
input_shape=input_shape)

```

```

    model = models.Sequential([
        base_model,
        layers.Flatten(),
        layers.Dense(512, activation='relu',
                      kernel_regularizer=regularizers.l2(0.001)),
        layers.Dropout(0.6),
        layers.Dense(1, activation='linear')
    ])

```

```

    model.compile(optimizer='adam', loss='mean_absolute_error',
metrics=['mae'])
    return model

def split_data(X, y, test_size=0.2, val_size=0.2):
    from sklearn.model_selection import train_test_split

    X_train_val, X_test, y_train_val, y_test = train_test_split(
        X, y, test_size=test_size, random_state=42
    )

    X_train, X_val, y_train, y_val = train_test_split(
        X_train_val, y_train_val, test_size=val_size, random_state=42
    )

    return X_train, X_val, X_test, y_train, y_val, y_test

def train_model(model, X_train, y_train, X_val, y_val,
data_augmentation=None):
    early_stopping = EarlyStopping(
        monitor='val_loss',
        patience=5,
        restore_best_weights=True
    )

    reduce_lr = ReduceLROnPlateau(
        monitor='val_loss',
        factor=0.2,
        patience=2,
        min_lr=1e-6
    )

    if data_augmentation:
        history = model.fit(
            data_augmentation.flow(X_train, y_train, batch_size=32),
            validation_data=(X_val, y_val),
            epochs=10,
            callbacks=[early_stopping, reduce_lr]
        )
    else:
        history = model.fit(
            X_train, y_train,
            validation_data=(X_val, y_val),
            epochs=50,
            batch_size=32,
            callbacks=[early_stopping, reduce_lr]
        )

```

```

    return history

def main3():
    zip_file_path = r"/content/drive/MyDrive/FCAI /DL_Assignment 1.zip"
    X, y = load_and_preprocess_data(zip_file_path)

    X_train, X_val, X_test, y_train, y_val, y_test = split_data(X, y)

    model = create_transfer_learning_model2()
    history = train_model(model, X_train, y_train, X_val, y_val)

    plt.figure(figsize=(10, 5))
    plt.plot(history.history['loss'], label='Training Loss')
    plt.plot(history.history['val_loss'], label='Validation Loss')
    plt.title('Model Training and Validation Loss')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.legend()
    plt.show()

    test_loss, test_mae = model.evaluate(X_test, y_test)
    print(f"Test Mean Absolute Error: {test_mae}")

    model.save("age_prediction_model.h5")
    print("Model saved as 'age_prediction_model.h5'.")

main3()

```

```

Downloading data from https://storage.googleapis.com/tensorflow/keras-
applications/resnet/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5
94765736/94765736 _____1s 0us/step
Epoch 1/50
412/412 _____114s 126ms/step - loss: 13.6702 - mae: 12.9076 -
val_loss: 20.1047 - val_mae: 19.6905 - learning_rate: 0.0010
Epoch 2/50
412/412 _____79s 52ms/step - loss: 9.4550 - mae: 9.1017 -
val_loss: 13.0077 - val_mae: 12.7903 - learning_rate: 0.0010
Epoch 3/50
412/412 _____42s 54ms/step - loss: 8.2578 - mae: 8.0673 -
val_loss: 8.6964 - val_mae: 8.5677 - learning_rate: 0.0010
Epoch 4/50
412/412 _____41s 53ms/step - loss: 7.9392 - mae: 7.8233 -
val_loss: 12.0236 - val_mae: 11.9364 - learning_rate: 0.0010
Epoch 5/50
412/412 _____41s 54ms/step - loss: 7.8534 - mae: 7.7737 -
val_loss: 10.2247 - val_mae: 10.1627 - learning_rate: 0.0010
Epoch 6/50
412/412 _____22s 54ms/step - loss: 6.8093 - mae: 6.7487 -
val_loss: 5.9902 - val_mae: 5.9335 - learning_rate: 2.0000e-04

```

Epoch 7/50  
412/412 ————— 40s 53ms/step - loss: 6.0977 - mae: 6.0420 -  
val\_loss: 6.0348 - val\_mae: 5.9826 - learning\_rate: 2.0000e-04

Epoch 8/50  
412/412 ————— 21s 52ms/step - loss: 5.9846 - mae: 5.9336 -  
val\_loss: 6.2288 - val\_mae: 6.1813 - learning\_rate: 2.0000e-04

Epoch 9/50  
412/412 ————— 22s 54ms/step - loss: 5.5749 - mae: 5.5276 -  
val\_loss: 5.7994 - val\_mae: 5.7528 - learning\_rate: 4.0000e-05

Epoch 10/50  
412/412 ————— 41s 53ms/step - loss: 5.3743 - mae: 5.3280 -  
val\_loss: 5.8085 - val\_mae: 5.7630 - learning\_rate: 4.0000e-05

Epoch 11/50  
412/412 ————— 41s 54ms/step - loss: 5.1918 - mae: 5.1466 -  
val\_loss: 5.8096 - val\_mae: 5.7654 - learning\_rate: 4.0000e-05

Epoch 12/50  
412/412 ————— 42s 56ms/step - loss: 5.1055 - mae: 5.0614 -  
val\_loss: 5.7732 - val\_mae: 5.7293 - learning\_rate: 8.0000e-06

Epoch 13/50  
412/412 ————— 41s 56ms/step - loss: 5.0229 - mae: 4.9791 -  
val\_loss: 5.7469 - val\_mae: 5.7034 - learning\_rate: 8.0000e-06

Epoch 14/50  
412/412 ————— 41s 56ms/step - loss: 5.0904 - mae: 5.0470 -  
val\_loss: 5.7407 - val\_mae: 5.6977 - learning\_rate: 8.0000e-06

Epoch 15/50  
412/412 ————— 41s 56ms/step - loss: 5.1569 - mae: 5.1140 -  
val\_loss: 5.7768 - val\_mae: 5.7343 - learning\_rate: 8.0000e-06

Epoch 16/50  
412/412 ————— 41s 56ms/step - loss: 4.9661 - mae: 4.9237 -  
val\_loss: 5.7677 - val\_mae: 5.7258 - learning\_rate: 8.0000e-06

Epoch 17/50  
412/412 ————— 41s 56ms/step - loss: 4.9136 - mae: 4.8717 -  
val\_loss: 5.7326 - val\_mae: 5.6908 - learning\_rate: 1.6000e-06


Epoch 18/50  
412/412 ————— 41s 56ms/step - loss: 4.9830 - mae: 4.9413 -  
val\_loss: 5.7548 - val\_mae: 5.7132 - learning\_rate: 1.6000e-06

Epoch 19/50  
412/412 ————— 21s 51ms/step - loss: 5.0009 - mae: 4.9594 -  
val\_loss: 5.7449 - val\_mae: 5.7035 - learning\_rate: 1.6000e-06

Epoch 20/50  
412/412 ————— 41s 52ms/step - loss: 4.9787 - mae: 4.9374 -  
val\_loss: 5.7442 - val\_mae: 5.7030 - learning\_rate: 1.0000e-06

Epoch 21/50  
412/412 ————— 43s 56ms/step - loss: 4.9214 - mae: 4.8802 -  
val\_loss: 5.7380 - val\_mae: 5.6969 - learning\_rate: 1.0000e-06

Epoch 22/50  
412/412 ————— 42s 59ms/step - loss: 4.8712 - mae: 4.8302 -  
val\_loss: 5.7373 - val\_mae: 5.6965 - learning\_rate: 1.0000e-06



129/129 ————— 4s 30ms/step - loss: 5.5172 - mae: 5.4754

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my\_model.keras')` or `keras.saving.save\_model(model, 'my\_model.keras')`.

Test Mean Absolute Error: 5.554823398590088  
Model saved as 'age\_prediction\_model.h5'.

## VGG16

```
import tensorflow as tf
from tensorflow.keras import layers, models, optimizers, regularizers
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
```

```
def load_and_preprocess_data(zip_file_path):
```

```
    images = []
    ages = []
```

```
    dataset_folder = '/content/drive/MyDrive/faces/utkcropped'
```

```
    for filename in os.listdir(dataset_folder):
        if filename.endswith('.jpg'):
```

```
            parts = filename.split('_')
```



```

    try:
        age = int(parts[0])

        if 10 <= age <= 90:
            img_path = os.path.join(dataset_folder, filename)
            img = cv2.imread(img_path)
            img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
            img = cv2.resize(img, (64, 64))

            images.append(img)
            ages.append(age)
    except ValueError:
        continue

X = np.array(images) / 255.0
y = np.array(ages)

return X, y

def create_transfer_learning_model1(input_shape=(64, 64, 3)):
    base_model = VGG16(weights='imagenet', include_top=False,
input_shape=input_shape)

    model = models.Sequential([
        base_model,
        layers.Flatten(),
        layers.Dense(256, activation='relu',
                      kernel_regularizer=regularizers.l2(0.001)),
        layers.Dropout(0.5),
        layers.Dense(1, activation='linear')
    ])

    model.compile(optimizer='adam', loss='mean_absolute_error',
metrics=['mae'])
    return model

def split_data(X, y, test_size=0.2, val_size=0.2):
    from sklearn.model_selection import train_test_split

    X_train_val, X_test, y_train_val, y_test = train_test_split(
        X, y, test_size=test_size, random_state=42
    )

    X_train, X_val, y_train, y_val = train_test_split(
        X_train_val, y_train_val, test_size=val_size, random_state=42
    )

```

```
return X_train, X_val, X_test, y_train, y_val, y_test
```

```
def train_model(model, X_train, y_train, X_val, y_val,  
data_augmentation=None):
```

```
    early_stopping = EarlyStopping(  
        monitor='val_loss',  
        patience=5,  
        restore_best_weights=True  
    )
```

```
    reduce_lr = ReduceLROnPlateau(  
        monitor='val_loss',  
        factor=0.2,  
        patience=2,  
        min_lr=1e-6  
    )
```

```
    if data_augmentation:  
        history = model.fit(  
            data_augmentation.flow(X_train, y_train, batch_size=32),  
            validation_data=(X_val, y_val),  
            epochs=10,  
            callbacks=[early_stopping, reduce_lr]  
        )  
    else:  
        history = model.fit(  
            X_train, y_train,  
            validation_data=(X_val, y_val),  
            epochs=50,  
            batch_size=32,  
            callbacks=[early_stopping, reduce_lr]  
        )
```

```
    return history
```

```
def main2():
```

```
    # Load Data
```

```
    zip_file_path = '/content/drive/MyDrive/faces/utkcropped'
```

```
    X, y = load_and_preprocess_data(zip_file_path)
```

```
    # Split Data
```

```
    X_train, X_val, X_test, y_train, y_val, y_test = split_data(X, y)
```

```
    # Create and Train Model
```

```
    model = create_transfer_learning_model1()
```

```
    history = train_model(model, X_train, y_train, X_val, y_val)
```

```

# Plot Training History
plt.figure(figsize=(10, 5))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Training and Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()

```

```

# Evaluate Model
test_loss, test_mae = model.evaluate(X_test, y_test)
print(f"Test Mean Absolute Error: {test_mae}")
model.save("age_prediction_model.h5")
print("Model saved as 'age_prediction_model.h5'")

```

main2()

```

Epoch 1/50
39/39 ————— 13s 186ms/step - loss: 26.2351 - mae: 25.8795 -
val_loss: 2.7915 - val_mae: 2.5894 - learning_rate: 0.0010
Epoch 2/50
39/39 ————— 3s 65ms/step - loss: 3.4029 - mae: 3.2133 -
val_loss: 2.7918 - val_mae: 2.6265 - learning_rate: 0.0010
Epoch 3/50
39/39 ————— 5s 67ms/step - loss: 3.4092 - mae: 3.2490 -
val_loss: 2.7893 - val_mae: 2.6427 - learning_rate: 0.0010
Epoch 4/50
39/39 ————— 3s 66ms/step - loss: 3.2943 - mae: 3.1513 -
val_loss: 2.7393 - val_mae: 2.6064 - learning_rate: 0.0010
Epoch 5/50
39/39 ————— 5s 66ms/step - loss: 3.3194 - mae: 3.1894 -
val_loss: 2.7340 - val_mae: 2.6119 - learning_rate: 0.0010
Epoch 6/50
39/39 ————— 5s 66ms/step - loss: 3.3952 - mae: 3.2754 -
val_loss: 2.8936 - val_mae: 2.7807 - learning_rate: 0.0010
Epoch 7/50
39/39 ————— 3s 64ms/step - loss: 3.2434 - mae: 3.1317 -
val_loss: 2.8970 - val_mae: 2.7908 - learning_rate: 0.0010
Epoch 8/50
39/39 ————— 3s 70ms/step - loss: 3.1580 - mae: 3.0519 -
val_loss: 2.7199 - val_mae: 2.6149 - learning_rate: 2.0000e-04
Epoch 9/50
39/39 ————— 3s 64ms/step - loss: 3.2489 - mae: 3.1442 -
val_loss: 2.6838 - val_mae: 2.5803 - learning_rate: 2.0000e-04
Epoch 10/50
39/39 ————— 3s 64ms/step - loss: 3.2023 - mae: 3.0991 -
val_loss: 2.7641 - val_mae: 2.6622 - learning_rate: 2.0000e-04
Epoch 11/50

```

39/39 ————— 2s 63ms/step - loss: 3.2589 - mae: 3.1572 -  
val\_loss: 2.7245 - val\_mae: 2.6240 - learning\_rate: 2.0000e-04  
Epoch 12/50  
39/39 ————— 3s 68ms/step - loss: 3.2312 - mae: 3.1307 -  
val\_loss: 2.6927 - val\_mae: 2.5923 - learning\_rate: 4.0000e-05  
Epoch 13/50  
39/39 ————— 3s 70ms/step - loss: 3.2412 - mae: 3.1408 -  
val\_loss: 2.6741 - val\_mae: 2.5741 - learning\_rate: 4.0000e-05  
Epoch 14/50  
39/39 ————— 5s 64ms/step - loss: 3.1609 - mae: 3.0610 -  
val\_loss: 2.6864 - val\_mae: 2.5867 - learning\_rate: 4.0000e-05  
Epoch 15/50  
39/39 ————— 2s 64ms/step - loss: 3.1312 - mae: 3.0316 -  
val\_loss: 2.6714 - val\_mae: 2.5720 - learning\_rate: 4.0000e-05  
Epoch 16/50  
39/39 ————— 3s 63ms/step - loss: 3.2141 - mae: 3.1149 -  
val\_loss: 2.7035 - val\_mae: 2.6045 - learning\_rate: 4.0000e-05  
Epoch 17/50  
39/39 ————— 3s 70ms/step - loss: 3.2147 - mae: 3.1158 -  
val\_loss: 2.6607 - val\_mae: 2.5621 - learning\_rate: 4.0000e-05  
Epoch 18/50  
39/39 ————— 3s 66ms/step - loss: 3.2288 - mae: 3.1303 -  
val\_loss: 2.7116 - val\_mae: 2.6134 - learning\_rate: 4.0000e-05  
Epoch 19/50  
39/39 ————— 3s 68ms/step - loss: 3.1037 - mae: 3.0055 -  
val\_loss: 2.6335 - val\_mae: 2.5356 - learning\_rate: 4.0000e-05  
Epoch 20/50  
39/39 ————— 2s 63ms/step - loss: 3.0945 - mae: 2.9967 -  
val\_loss: 2.6687 - val\_mae: 2.5712 - learning\_rate: 4.0000e-05  
Epoch 21/50  
39/39 ————— 3s 64ms/step - loss: 2.9893 - mae: 2.8920 -  
val\_loss: 2.6039 - val\_mae: 2.5069 - learning\_rate: 4.0000e-05  
Epoch 22/50  
39/39 ————— 3s 68ms/step - loss: 3.0786 - mae: 2.9817 -  
val\_loss: 2.6558 - val\_mae: 2.5592 - learning\_rate: 4.0000e-05  
Epoch 23/50  
39/39 ————— 3s 65ms/step - loss: 2.9056 - mae: 2.8091 -  
val\_loss: 2.6050 - val\_mae: 2.5088 - learning\_rate: 4.0000e-05  
Epoch 24/50  
39/39 ————— 3s 64ms/step - loss: 2.9494 - mae: 2.8533 -  
val\_loss: 2.5791 - val\_mae: 2.4831 - learning\_rate: 8.0000e-06  
Epoch 25/50  
39/39 ————— 2s 63ms/step - loss: 2.8896 - mae: 2.7935 -  
val\_loss: 2.5222 - val\_mae: 2.4262 - learning\_rate: 8.0000e-06  
Epoch 26/50  
39/39 ————— 2s 62ms/step - loss: 2.8007 - mae: 2.7047 -  
val\_loss: 2.5254 - val\_mae: 2.4295 - learning\_rate: 8.0000e-06  
Epoch 27/50  
39/39 ————— 2s 64ms/step - loss: 2.8189 - mae: 2.7230 -  
val\_loss: 2.5047 - val\_mae: 2.4089 - learning\_rate: 8.0000e-06

Epoch 28/50  
39/39 ————— 3s 68ms/step - loss: 2.9107 - mae: 2.8149 -  
val\_loss: 2.5131 - val\_mae: 2.4174 - learning\_rate: 8.0000e-06

Epoch 29/50  
39/39 ————— 3s 66ms/step - loss: 2.8819 - mae: 2.7861 -  
val\_loss: 2.4885 - val\_mae: 2.3929 - learning\_rate: 8.0000e-06

Epoch 30/50  
39/39 ————— 5s 64ms/step - loss: 2.8542 - mae: 2.7586 -  
val\_loss: 2.5031 - val\_mae: 2.4076 - learning\_rate: 8.0000e-06

Epoch 31/50  
39/39 ————— 2s 64ms/step - loss: 2.8179 - mae: 2.7224 -  
val\_loss: 2.4876 - val\_mae: 2.3921 - learning\_rate: 8.0000e-06

Epoch 32/50  
39/39 ————— 2s 62ms/step - loss: 2.7453 - mae: 2.6499 -  
val\_loss: 2.4991 - val\_mae: 2.4038 - learning\_rate: 8.0000e-06

Epoch 33/50  
39/39 ————— 3s 68ms/step - loss: 2.8924 - mae: 2.7971 -  
val\_loss: 2.5676 - val\_mae: 2.4724 - learning\_rate: 8.0000e-06

Epoch 34/50  
39/39 ————— 5s 64ms/step - loss: 2.9100 - mae: 2.8148 -  
val\_loss: 2.4914 - val\_mae: 2.3962 - learning\_rate: 1.6000e-06

Epoch 35/50  
39/39 ————— 5s 65ms/step - loss: 2.8239 - mae: 2.7287 -  
val\_loss: 2.4810 - val\_mae: 2.3858 - learning\_rate: 1.6000e-06

Epoch 36/50  
39/39 ————— 3s 69ms/step - loss: 2.8243 - mae: 2.7291 -  
val\_loss: 2.4809 - val\_mae: 2.3857 - learning\_rate: 1.6000e-06


Epoch 37/50  
39/39 ————— 3s 66ms/step - loss: 3.0190 - mae: 2.9238 -  
val\_loss: 2.5007 - val\_mae: 2.4055 - learning\_rate: 1.6000e-06

Epoch 38/50  
39/39 ————— 5s 64ms/step - loss: 2.8173 - mae: 2.7222 -  
val\_loss: 2.4961 - val\_mae: 2.4010 - learning\_rate: 1.6000e-06

Epoch 39/50  
39/39 ————— 5s 64ms/step - loss: 2.6913 - mae: 2.5962 -  
val\_loss: 2.4949 - val\_mae: 2.3998 - learning\_rate: 1.0000e-06

Epoch 40/50  
39/39 ————— 5s 64ms/step - loss: 2.9207 - mae: 2.8256 -  
val\_loss: 2.4920 - val\_mae: 2.3969 - learning\_rate: 1.0000e-06

Epoch 41/50  
39/39 ————— 2s 63ms/step - loss: 2.9778 - mae: 2.8827 -  
val\_loss: 2.4854 - val\_mae: 2.3903 - learning\_rate: 1.0000e-06



13/13 ————— 0s 38ms/step - loss: 2.4259 - mae: 2.3308

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my\_model.keras')` or `keras.saving.save\_model(model, 'my\_model.keras')`.

Test Mean Absolute Error: 2.340360403060913  
Model saved as 'age\_prediction\_model.h5'

## 8. Select the best model, use it to make age predictions on face images of your team, and plot these images with the estimated ages.

```
import os
import cv2
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt

def predict_and_visualize(model_path, image_paths):

    model = tf.keras.models.load_model(model_path)

    images = []
    valid_image_paths = []

    for img_path in image_paths:
        if not os.path.exists(img_path):
            print(f"Image not found: {img_path}")
```

```

        continue
    img = cv2.imread(img_path)
    if img is None:
        print(f"Failed to load image: {img_path}")
        continue
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    img = cv2.resize(img, (64, 64))
    images.append(img)
    valid_image_paths.append(img_path)

if not images:
    print("No valid images to process.")
    return

images = np.array(images) / 255.0

predictions = model.predict(images)
predicted_ages = predictions.flatten()

plt.figure(figsize=(12, 6))
for i, img in enumerate(images):
    plt.subplot(1, len(images), i + 1)
    plt.imshow(img)
    plt.title(f"Predicted Age: {int(predicted_ages[i])}")
    plt.axis('off')
plt.show()

image_paths = [
    "/content/Menna.jpg",
    "/content/Mariam.jpg",
    "/content/jana.jpg",
    "/content/Zaynab.jpg"
]

```

```
predict_and_visualize("age_prediction_model.h5", image_paths)
```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile\_metrics` will be empty until you train or evaluate the model.

1/1 ————— 0s 370ms/step



## VGG16

```
import tensorflow as tf
from tensorflow.keras import layers, models, optimizers, regularizers
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
import os # Import the os module
from tensorflow.keras.applications import VGG16, ResNet50
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import zipfile
import cv2
```

*# Deep Learning Libraries*

```
import tensorflow as tf
from tensorflow.keras import layers, models, optimizers, regularizers
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.applications import VGG16, ResNet50
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
def load_and_preprocess_data(zip_file_path):
```

```
    images = []
    ages = []
```

```
    dataset_folder = '/content/drive/MyDrive/FCAI /utkcropped'
```

```
    for filename in os.listdir(dataset_folder):
        if filename.endswith('.jpg'):
```

```
            parts = filename.split('_')
            try:
                age = int(parts[0])
```

```
                if 10 <= age <= 90:
                    img_path = os.path.join(dataset_folder, filename)
                    img = cv2.imread(img_path)
                    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
                    img = cv2.resize(img, (64, 64))
```

```
                    images.append(img)
                    ages.append(age)
```

```
            except ValueError:
                continue
```

```
    X = np.array(images) / 255.0
```



```

y = np.array(ages)

return X, y

def create_transfer_learning_model1(input_shape=(64, 64, 3)):
    base_model = VGG16(weights='imagenet', include_top=False,
input_shape=input_shape)

    model = models.Sequential([
        base_model,
        layers.Flatten(),
        layers.Dense(256, activation='relu',
                      kernel_regularizer=regularizers.l2(0.001)),
        layers.Dropout(0.5),
        layers.Dense(1, activation='linear')
    ])

    model.compile(optimizer='adam', loss='mean_absolute_error',
metrics=['mae'])
    return model

def split_data(X, y, test_size=0.2, val_size=0.2):
    from sklearn.model_selection import train_test_split

    X_train_val, X_test, y_train_val, y_test = train_test_split(
        X, y, test_size=test_size, random_state=42
    )

    X_train, X_val, y_train, y_val = train_test_split(
        X_train_val, y_train_val, test_size=val_size, random_state=42
    )

    return X_train, X_val, X_test, y_train, y_val, y_test

def train_model(model, X_train, y_train, X_val, y_val,
data_augmentation=None):

    early_stopping = EarlyStopping(
        monitor='val_loss',
        patience=5,
        restore_best_weights=True
    )

    reduce_lr = ReduceLROnPlateau(
        monitor='val_loss',
        factor=0.2,

```

```

        patience=2,
        min_lr=1e-6
    )

    if data_augmentation:
        history = model.fit(
            data_augmentation.flow(X_train, y_train, batch_size=32),
            validation_data=(X_val, y_val),
            epochs=10,
            callbacks=[early_stopping, reduce_lr]
        )
    else:
        history = model.fit(
            X_train, y_train,
            validation_data=(X_val, y_val),
            epochs=50,
            batch_size=32,
            callbacks=[early_stopping, reduce_lr]
        )

    return history

import cv2 # Import the OpenCV Library
import numpy as np #Import the numpy Library

def main2():
    # Load Data
    zip_file_path = '/content/drive/MyDrive/FCAI /DL_Assignment 1.zip'
    X, y = load_and_preprocess_data(zip_file_path)

    # Split Data
    X_train, X_val, X_test, y_train, y_val, y_test = split_data(X, y)

    # Create and Train Model
    model = create_transfer_learning_model1()
    history = train_model(model, X_train, y_train, X_val, y_val)

    # Plot Training History
    plt.figure(figsize=(10, 5))
    plt.plot(history.history['loss'], label='Training Loss')
    plt.plot(history.history['val_loss'], label='Validation Loss')
    plt.title('Model Training and Validation Loss')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.legend()
    plt.show()

    # Evaluate Model
    test_loss, test_mae = model.evaluate(X_test, y_test)
    print(f"Test Mean Absolute Error: {test_mae}")

```

```
model.save("age_prediction_model.h5")
print("Model saved as 'age_prediction_model.h5'")
```

```
main2()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-
applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58889256/58889256 _____0s 0us/step
```

```
Epoch 1/50
```

```
412/412 _____54s 98ms/step - loss: 24.1014 - mae: 23.8853 -
val_loss: 13.4696 - val_mae: 13.3534 - learning_rate: 0.0010
```

```
Epoch 2/50
```

```
412/412 _____59s 64ms/step - loss: 13.3339 - mae: 13.2258 -
val_loss: 13.4993 - val_mae: 13.4077 - learning_rate: 0.0010
```

```
Epoch 3/50
```

```
412/412 _____28s 67ms/step - loss: 13.3378 - mae: 13.2460 -
val_loss: 12.7734 - val_mae: 12.6873 - learning_rate: 0.0010
```

```
Epoch 4/50
```

```
412/412 _____41s 68ms/step - loss: 12.3991 - mae: 12.3128 -
val_loss: 10.1688 - val_mae: 10.0839 - learning_rate: 0.0010
```

```
Epoch 5/50
```

```
412/412 _____40s 65ms/step - loss: 10.5264 - mae: 10.4438 -
val_loss: 9.8306 - val_mae: 9.7496 - learning_rate: 0.0010
```

```
Epoch 6/50
```

```
412/412 _____42s 68ms/step - loss: 9.4403 - mae: 9.3630 -
val_loss: 8.5121 - val_mae: 8.4421 - learning_rate: 0.0010
```

```
Epoch 7/50
```

```
412/412 _____27s 65ms/step - loss: 8.6056 - mae: 8.5376 -
val_loss: 8.2044 - val_mae: 8.1414 - learning_rate: 0.0010
```

```
Epoch 8/50
```

```
412/412 _____41s 65ms/step - loss: 8.2628 - mae: 8.2011 -
val_loss: 7.6528 - val_mae: 7.5935 - learning_rate: 0.0010
```

```
Epoch 9/50
```

```
412/412 _____27s 65ms/step - loss: 7.9129 - mae: 7.8542 -
val_loss: 7.7065 - val_mae: 7.6521 - learning_rate: 0.0010
```

```
Epoch 10/50
```

```
412/412 _____41s 64ms/step - loss: 7.7057 - mae: 7.6531 -
val_loss: 7.7257 - val_mae: 7.6759 - learning_rate: 0.0010
```

```
Epoch 11/50
```

```
412/412 _____27s 65ms/step - loss: 7.1314 - mae: 7.0817 -
val_loss: 7.1162 - val_mae: 7.0680 - learning_rate: 2.0000e-04
```

```
Epoch 12/50
```

```
412/412 _____42s 67ms/step - loss: 6.5581 - mae: 6.5101 -
val_loss: 6.8062 - val_mae: 6.7596 - learning_rate: 2.0000e-04
```


```
Epoch 13/50
```

```
412/412 _____41s 67ms/step - loss: 6.2543 - mae: 6.2077 -
val_loss: 6.7856 - val_mae: 6.7403 - learning_rate: 2.0000e-04
```

```
Epoch 14/50
```

```
412/412 _____28s 67ms/step - loss: 6.0338 - mae: 5.9887 -
val_loss: 6.6632 - val_mae: 6.6192 - learning_rate: 2.0000e-04
```

Epoch 15/50  
412/412 —————40s 65ms/step - loss: 5.9056 - mae: 5.8620 -  
val\_loss: 6.6917 - val\_mae: 6.6493 - learning\_rate: 2.0000e-04  
Epoch 16/50  
412/412 —————27s 65ms/step - loss: 5.6445 - mae: 5.6022 -  
val\_loss: 6.6551 - val\_mae: 6.6138 - learning\_rate: 2.0000e-04  
Epoch 17/50  
412/412 —————27s 65ms/step - loss: 5.4421 - mae: 5.4012 -  
val\_loss: 6.5996 - val\_mae: 6.5601 - learning\_rate: 2.0000e-04  
Epoch 18/50  
412/412 —————28s 68ms/step - loss: 5.3080 - mae: 5.2690 -  
val\_loss: 6.7268 - val\_mae: 6.6894 - learning\_rate: 2.0000e-04  
Epoch 19/50  
412/412 —————28s 68ms/step - loss: 5.0909 - mae: 5.0538 -  
val\_loss: 6.6902 - val\_mae: 6.6546 - learning\_rate: 2.0000e-04  
Epoch 20/50  
412/412 —————28s 68ms/step - loss: 4.8938 - mae: 4.8581 -  
val\_loss: 6.5635 - val\_mae: 6.5281 - learning\_rate: 4.0000e-05  
Epoch 21/50  
412/412 —————28s 69ms/step - loss: 4.6128 - mae: 4.5774 -  
val\_loss: 6.5574 - val\_mae: 6.5222 - learning\_rate: 4.0000e-05  
Epoch 22/50  
412/412 —————41s 70ms/step - loss: 4.4972 - mae: 4.4621 -  
val\_loss: 6.5963 - val\_mae: 6.5617 - learning\_rate: 4.0000e-05  
Epoch 23/50  
412/412 —————40s 68ms/step - loss: 4.4202 - mae: 4.3856 -  
val\_loss: 6.5977 - val\_mae: 6.5637 - learning\_rate: 4.0000e-05  
Epoch 24/50  
412/412 —————42s 70ms/step - loss: 4.3193 - mae: 4.2853 -  
val\_loss: 6.5717 - val\_mae: 6.5378 - learning\_rate: 8.0000e-06  
Epoch 25/50  
412/412 —————42s 72ms/step - loss: 4.3513 - mae: 4.3174 -  
val\_loss: 6.5693 - val\_mae: 6.5355 - learning\_rate: 8.0000e-06  
Epoch 26/50  
412/412 —————30s 72ms/step - loss: 4.3029 - mae: 4.2692 -  
val\_loss: 6.5666 - val\_mae: 6.5328 - learning\_rate: 1.6000e-06



129/129 ————— 5s 37ms/step - loss: 6.3550 - mae: 6.3198

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my\_model.keras')` or `keras.saving.save\_model(model, 'my\_model.keras')`.

Test Mean Absolute Error: 6.368239879608154  
Model saved as 'age\_prediction\_model.h5'

```
import os
import cv2
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt

def predict_and_visualize(model_path, image_paths):

    model = tf.keras.models.load_model(model_path)

    images = []
    valid_image_paths = []

    for img_path in image_paths:
        if not os.path.exists(img_path):
            print(f"Image not found: {img_path}")
            continue
        img = cv2.imread(img_path)
        if img is None:
            print(f"Failed to load image: {img_path}")
```

```

        continue
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    img = cv2.resize(img, (64, 64))
    images.append(img)
    valid_image_paths.append(img_path)

if not images:
    print("No valid images to process.")
    return

images = np.array(images) / 255.0

predictions = model.predict(images)
predicted_ages = predictions.flatten()

plt.figure(figsize=(12, 6))
for i, img in enumerate(images):
    plt.subplot(1, len(images), i + 1)
    plt.imshow(img)
    plt.title(f"Predicted Age: {int(predicted_ages[i])}")
    plt.axis('off')
plt.show()

image_paths = [
    "/content/Menna.jpg",
    "/content/Mariam.jpg",
    "/content/Jana.jpg",
    "/content/Zaynab.jpg"
]

predict_and_visualize("age_prediction_model.h5", image_paths)

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to
be built. `model.compile_metrics` will be empty until you train or evaluate
the model.

```

1/1 ————— 2s 2s/step



Model	Mean Absolute Error
Custom CNN model 1	6.466558933258057
customCNNmodel2	6.37539005279541
ResNet50	5.554823398590088
VGG16 1	2.340360403060913
VGG16 2	6.3198

**Conclusion:**

**As seen in the previous table, VGG16 version 1 got the lowest MAE.**